

Using MongoDB at x.ai




x.ai a personal assistant
who schedules meetings for you


matt casey
@xdotai

Magically Schedule Meetings

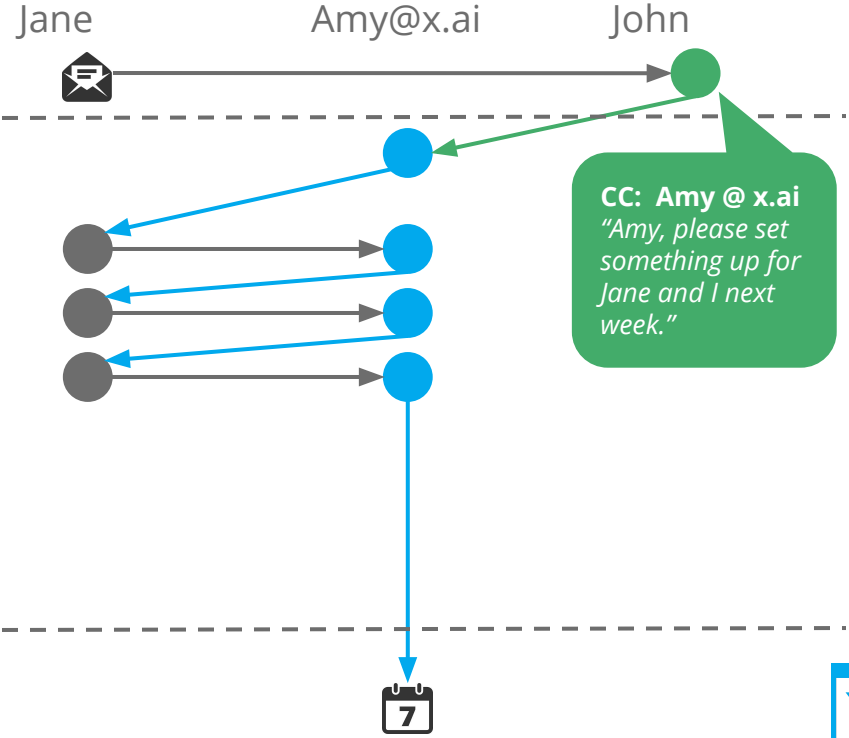
Just Cc: Amy@x.ai



Pain



Solution



System characteristics

- Need quick response
- Learning algorithms require large training data set
- # meetings scale linearly with # users
- 1 user meets with N people
- People share meeting, places, times and company
- Social relationships (assistants, coordinators)



Technical challenges


- Natural language understanding with extremely high accuracy
- Natural conversation over email with people
- Complex data relationship
- Optimize for sparse data
- Speed of development and change



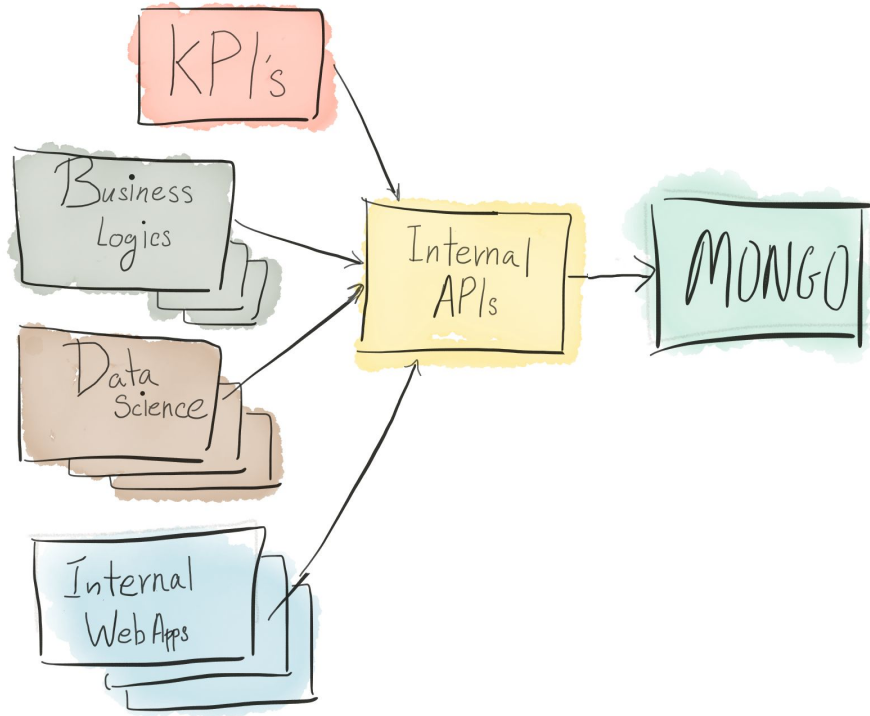
Stack



Queue Based Architecture



Data Access



- Schema Standardization
- Mongoose
- Supported by Elastic Load Balancer



Data Models: Let's Get Started

Here's a few tips:

- Mongoose.js ODM for CRUD services
- mongoose-express-restify to provide a REST API
- MongoDB MMS for monitoring and backups
- Keep old data up-to-date with schema changes
- Don't over-design, use models that are easy to change



How Do We Model Time?

Use Case:

- Capture times and constraints from an email for scheduling a meeting
- Data is created, saved, and read by a human

“Let’s meet the first week of February next year”

```
time: {  
    start: new Date(2017, 1, 1),  
    end: new Date(2017, 1, 2)  
}
```



Try Again: Our Second Model

Use Case:

- A nested model captures the original intent
- Makes it possible for machines to detect and respond

```
timeV2: {  
  within: {  
    reference: {  
      weekOfYear: 4,  
      year: 2017  
    }  
  }  
}
```



Still, some room for improvement

- Storing data closer to text makes it easier to improve accuracy
- But is harder to interpret and validate without tools

```
timeV3: [  
  { timeId: id0, text: 'first week', weekOfMonth: 1 },  
  { timeId: id0, text: 'February', month: 2 },  
  { timeId: id0, text: 'next year', year: 2017 }  
]
```



Edge cases dealing with time

- Scheduling phone calls for people in different timezones
- Dealing with partial information
 - ***“5pm is good” or “next week”***
- Irrelevant times
 - ***“Can’t wait to see you next month. Can we talk tomorrow?”***
- Need context for timezones: natural language vs tz database
 - Identify: ***“IST,” “CST,” “EST”***
 - ***“I’m -5 from Dave” or “3pm London time”***
 - Etc/GMT+3 is not what you think

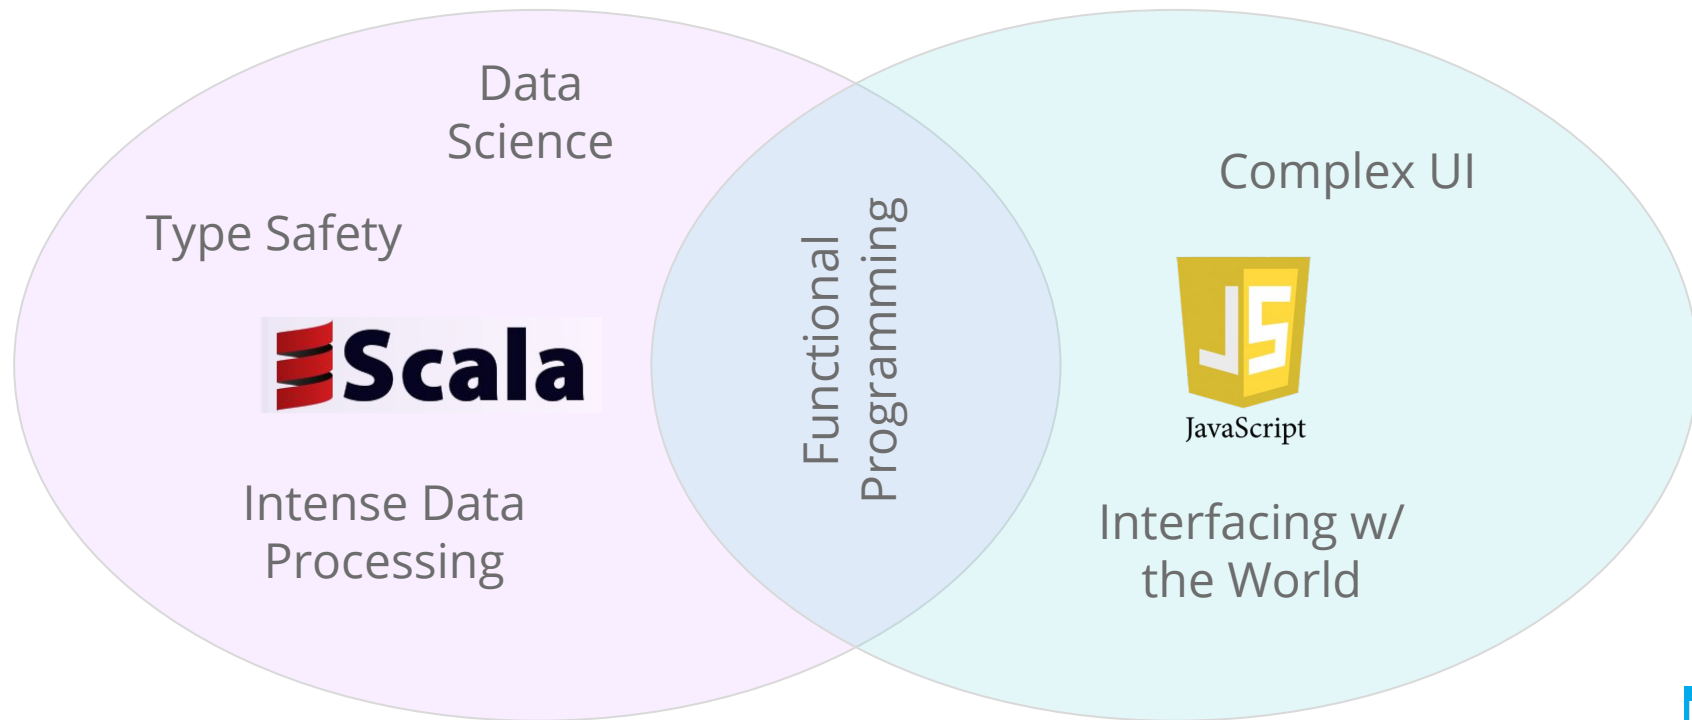


Going from human to machine

- Machine detects and predicts predicts times
- Human work shifts from feeding new information to correcting and re-affirming
- Focus shifts to generating more information and ensuring product quality



Scala & JS - Happily Ever After



Mongoose.js

object modeling and querying library for mongodb with node.js.



- ODM
- Easy to set up, convenient
- Schema enforcement
- Use with caution: getters, setters, virtual fields
- Client side joins across collections and databases
- Nested documents with automatic `_id` reference



- Applies to only one language
- ``populate`` mutates the object
- Functional Programming expects domain “events” vs objects
- Impossible to ‘clone’ instances
- using schema requires db connection
- No read validation - errors on updates



Mongo & Scala

Mongo

```
{
  "messageId" : "<somerandommessage@xxxx>",
  "classifierLabels" : [
    {
      "classifier" : "MeetingClassifier",
      "class" : "A"
    }
  ],
  "featureVector" : [
    {
      "featureType" : "TaggedOneGram",
      "featureKeys" : [
        {
          "featureKey" : "work",
          "featureValue" : 3
        },
        {
          "featureKey" : "that",
          "featureValue" : 1
        },
        ...
      ]
    },
    ...
  ]
}
```



Scala

```
case class Features (
  messageId: String,
  classes: List[Classes],
  featureVector: List[FeatureVectors]
)

case class Classes (
  classifier: String,
  class: String
)

case class FeatureVectors(
  featureType: String,
  featureKeys: List[FeatureKeys]
)

case class FeatureKey(
  featureKey: String,
  featureValue: 1
)
```



Schemaless vs. Typesafe

- Discrepancy between models in Scala versus stored data in mongo.
 - MongoDB stores free form nested structures
 - Scala relies on strict, type safed models for data
- Models in flux
 - Same format for both data extraction and decision making layers
 - We are continuously learning new edge case about scheduling meetings



Updating schemas in production

Step #1: Design new schema

- Avoid mutating the meaning of a field, add a new one instead
`timeEntities => timeEntitiesV2`
- Don't stop writing to the old field until other processes are updated

Step #2: Migrate old documents, if possible, or use a 'blacklisted' field for training

Step #3: Transition all processes that read from schema to look for the new model

Design Tips:

- Avoid saving state to the db, keep the data as close to reality as possible.
- Separate collections for each service
- Future-proof your model, a list of objects is more extendable than a list of strings:
`actions: ['SENT', 'DELIVERED']`
`actions: [{ timestamp, action: 'SENT' }, { timestamp, action: 'DELIVERED' }]`



Lessons Learned

- Let the data tell the story, keep track of what was said or happened
- Avoid foreign keys to “moving” documents
- Be cautious of too much buy-in to your tools
- Write tests from Day One
- Make backups!



matt @ human.x.ai

cto & founder

25 Broadway. 9th Floor
New York, 10004 NY

E: hello@human.x.ai

T: @xdotai

Visit x.ai to join the waitlist



Embedding vs. Referencing

Embedding

```
{
  host :
  {
    name : {.....},
    nicknames : [String],
    phones : [{Type: String}]
    primaryEmail : String,
    secondaryEmails : [String],
    title : String,
    signatures: [String],
    .....
  },
  travelTime : String,
  status : String,
  timezone : String,
  duration : Number,
  .....
}
```

Referencing

```
{
  host : Participant,
  travelTime : String,
  status : String,
  timezone : String,
  duration : Number,
  .....
}

Participant
{
  name : {.....},
  nicknames : [String],
  phones : [{Type: String}]
  primaryEmail : String,
  secondaryEmails : [String],
  title : String,
  signatures: [String],
  .....
},
```

Considerations

- Query patterns
- Access to embedded doc
- # references to a doc
- Application level join
- 1-way or 2-way referencing



Unsolved Schema issues

- Keeping schema changes in sync with Mongoose and Scala
- No libraries or tests to make sure the contracts don't break
 - ProtoBuffers or Avro?
 - discriminator key
- Managing a 'context collection' where multiple copies of production are stored per email
- Deploying schema changes sometimes requires re-training a model
- Schema validation from Scala supercedes Mongoose

